



**Министерство науки и высшего образования Российской Федерации
Федеральное государственное бюджетное образовательное учреждение
высшего образования
«Московский государственный технический университет
имени Н.Э. Баумана
(национальный исследовательский университет)»
(МГТУ им. Н.Э. Баумана)**

**Факультет «Информатика и системы управления»
Кафедра ИУ5 «Системы обработки информации и управления»**

**Отчет по лабораторной работы №7
по дисциплине «Методы машинного обучения»
по теме «Алгоритмы Actor-Critic»**

**Выполнил:
студент группы № ИУ5-21М
Торжков М.С.
подпись, дата**

**Проверила:
Балашов А.М.
подпись, дата**

2023 г.

Задание.

Реализуйте любой алгоритм семейства Actor-Critic для произвольной среды

```
! pip install gymnasium
import gymnasium as gym
import numpy as np
from itertools import count
from collections import namedtuple
```

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.distributions import Categorical
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting gymnasium

Downloading gymnasium-0.28.1-py3-none-any.whl (925 kB)

0:00:00 925.5/925.5 kB 19.6 MB/s eta

Requirement already satisfied: numpy>=1.21.0 in
/usr/local/lib/python3.10/dist-packages (from gymnasium) (1.22.4)
Collecting jax-jumpy>=1.0.0 (from gymnasium)

Downloading jax_jumpy-1.0.0-py3-none-any.whl (20 kB)

Requirement already satisfied: cloudpickle>=1.2.0 in
/usr/local/lib/python3.10/dist-packages (from gymnasium) (2.2.1)

Requirement already satisfied: typing-extensions>=4.3.0 in
/usr/local/lib/python3.10/dist-packages (from gymnasium) (4.5.0)

Collecting farama-notifications>=0.0.1 (from gymnasium)

Downloading Farama_Notifications-0.0.4-py3-none-any.whl (2.5 kB)

Installing collected packages: farama-notifications, jax-jumpy,
gymnasium

Successfully installed farama-notifications-0.0.4 gymnasium-0.28.1
jax-jumpy-1.0.0

```
!pip install pygame
```

```
import os
os.environ['SDL_VIDEODRIVER'] = 'dummy'
import pygame
pygame.display.set_mode((640, 480))
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Requirement already satisfied: pygame in
/usr/local/lib/python3.10/dist-packages (2.3.0)

<Surface(640x480x32 SW)>

```
# Cart Pole
```

```
CONST_ENV_NAME = 'Acrobot-v1'
env = gym.make(CONST_ENV_NAME)
```

```

GAMMA = 0.99
SavedAction = namedtuple('SavedAction', ['log_prob', 'value'])

class Policy(nn.Module):
    def __init__(self):
        super(Policy, self).__init__()
        self.affine1 = nn.Linear(6, 128)

        # actor's layer
        self.action_head = nn.Linear(128, 3)

        # critic's layer
        self.value_head = nn.Linear(128, 1)

        # action & reward buffer
        self.saved_actions = []
        self.rewards = []

    def forward(self, x):
        x = F.relu(self.affine1(x))

        # actor: choses action to take from state s_t
        # by returning probability of each action
        action_prob = F.softmax(self.action_head(x), dim=-1)

        # critic: evaluates being in the state s_t
        state_values = self.value_head(x)

        # return values for both actor and critic as a tuple of 2 values:
        # 1. a list with the probability of each action over the action
        # 2. the value from state s_t
        return action_prob, state_values

model = Policy()
optimizer = optim.AdamW(model.parameters(), lr=1e-3)
eps = np.finfo(np.float32).eps.item()

def select_action(state):
    state = torch.from_numpy(state).float()
    probs, state_value = model(state)

    # create a categorical distribution over the list of probabilities
    # of actions
    m = Categorical(probs)

    # and sample an action using the distribution
    action = m.sample()

    # save to action buffer

```

```

    model.saved_actions.append(SavedAction(m.log_prob(action),
state_value))

    # the action to take (left or right)
    return action.item()

def finish_episode():
    """
    Training code. Calculates actor and critic loss and performs
    backprop.
    """
    R = 0
    saved_actions = model.saved_actions
    policy_losses = [] # list to save actor (policy) loss
    value_losses = [] # list to save critic (value) loss
    returns = [] # list to save the true values

    # calculate the true value using rewards returned from the
environment
    for r in model.rewards[::1]:
        # calculate the discounted value
        R = r + GAMMA * R
        returns.insert(0, R)

    returns = torch.tensor(returns)
    returns = (returns - returns.mean()) / (returns.std() + eps)

    for (log_prob, value), R in zip(saved_actions, returns):
        advantage = R - value.item()

        # calculate actor (policy) loss
        policy_losses.append(-log_prob * advantage)

        # calculate critic (value) loss using L1 smooth loss
        value_losses.append(F.smooth_l1_loss(value, torch.tensor([R])))

    # reset gradients
    optimizer.zero_grad()

    # sum up all the values of policy_losses and value_losses
    loss = torch.stack(policy_losses).sum() +
torch.stack(value_losses).sum()

    # perform backprop
    loss.backward()
    optimizer.step()

    # reset rewards and action buffer

```

```

del model.rewards[:]
del model.saved_actions[:]

running_reward = -500

# run infinitely many episodes
for i_episode in count(1):
    #print(running_reward)
    # reset environment and episode reward
    state, _ = env.reset()
    ep_reward = 0
    # for each episode, only run 9999 steps so that we don't
    # infinite loop while learning
    for t in range(1, 9999):
        # select action from policy
        action = select_action(state)
        # take the action
        state, reward, done, truncated, _ = env.step(action)
        model.rewards.append(reward)
        ep_reward += reward
        if done or truncated:
            break
    print(ep_reward)
    # update cumulative reward
    running_reward = 0.05 * ep_reward + (1 - 0.05) * running_reward
    # perform backprop
    finish_episode()
    # log results
    if i_episode % 10 == 0:
        print(f"Episode {i_episode}\tLast reward: {ep_reward:.2f}\tAverage
reward: {running_reward:.2f}")
        # check if we have "solved" the cart pole problem
        if running_reward > env.spec.reward_threshold*2:
            print(f"Solved! Running reward is now {running_reward} and the
last episode runs to {t} time steps!")
            break
    env2 = gym.make(CONST_ENV_NAME, render_mode='human')
    # reset environment and episode reward
    state, _ = env2.reset()
    ep_reward = 0
    # for each episode, only run 9999 steps so that we don't
    # infinite loop while learning
    for t in range(1, 10000):
        # select action from policy
        action = select_action(state)
        # take the action
        state, reward, done, _, _ = env2.step(action)
        model.rewards.append(reward)
        ep_reward += reward

```

```
if done:  
    break
```

```
-500.0  
-500.0  
-500.0  
-500.0  
-500.0  
-500.0  
-500.0  
-500.0  
-461.0  
-500.0
```

Episode 10 Last reward: -500.00 Average reward: -498.15

```
-500.0  
-500.0  
-500.0  
-500.0  
-500.0  
-500.0  
-500.0  
-500.0  
-479.0  
-500.0  
-500.0
```

Episode 20 Last reward: -500.00 Average reward: -497.94

```
-340.0  
-500.0  
-500.0  
-425.0  
-400.0  
-448.0  
-500.0  
-500.0  
-500.0  
-500.0
```

Episode 30 Last reward: -500.00 Average reward: -484.98

```
-478.0  
-471.0  
-499.0  
-500.0  
-408.0  
-472.0  
-427.0  
-483.0  
-500.0  
-500.0
```

Episode 40 Last reward: -500.00 Average reward: -480.72

```
-500.0  
-361.0  
-500.0
```

-500.0

-500.0

-333.0

-340.0

-299.0

-342.0

-377.0

Episode 50 Last reward: -377.00 Average reward: -447.46

-347.0

-482.0

-344.0

-259.0

-311.0

-304.0

-500.0

-326.0

-263.0

-295.0

Episode 60 Last reward: -295.00 Average reward: -404.17

-184.0

-201.0

-328.0

-300.0

-199.0

-410.0

-288.0

-368.0

-339.0

-300.0

Episode 70 Last reward: -300.00 Average reward: -361.39

-292.0

-274.0

-332.0

-312.0

-233.0

-222.0

-363.0

-238.0

-285.0

-252.0

Episode 80 Last reward: -252.00 Average reward: -328.22

-226.0

-259.0

-354.0

-217.0

-500.0

-197.0

-181.0

-251.0

-198.0

-220.0
Episode 90 Last reward: -220.00 Average reward: -299.42
-244.0
-195.0
-233.0
-165.0
-191.0
-170.0
-202.0
-229.0
-218.0
-187.0
Episode 100 Last reward: -187.00 Average reward: -260.65
-166.0
-153.0
-203.0
-120.0
-218.0
-176.0
-249.0
-171.0
-255.0
-131.0
Episode 110 Last reward: -131.00 Average reward: -230.52
-226.0
-222.0
-176.0
-182.0
-137.0
-158.0
-197.0
-217.0
-185.0
-154.0
Episode 120 Last reward: -154.00 Average reward: -211.80
-183.0
-173.0
-153.0
-153.0
-116.0
Solved! Running reward is now -198.72727598097 and the last episode runs to 117 time steps!